

CSC 212: Data Structures and Abstractions

Balanced trees (part 2)

Prof. Marco Alvarez

Department of Computer Science and Statistics
University of Rhode Island

Fall 2025



BST Rotations


- A **rotation** is a O(1)-time local operation that preserves the **BST order property** while changing the tree's structure

Right rotation at node y

- ✓ requires y's left child x to be *non-null*
- ✓ elevates x to become the subtree root
- ✓ y becomes x's right child
- ✓ x's original right child becomes y's left child

Left rotation at node x

- ✓ requires x's right child y to be *non-null*
- ✓ elevates y to become the subtree root
- ✓ x becomes y's left child
- ✓ y's original left child becomes x's right child




3

Rotations

Example: left rotation

left-rotate(x)




<https://www.formosa1544.com/2021/04/30/build-the-forest-in-python-series-red-black-tree/>

4

Example: right rotation

right-rotate(x)




<https://www.formosa1544.com/2021/04/30/build-the-forest-in-python-series-red-black-tree/>

5

Practice

• Perform the following operations in sequence

- ✓ rotate-left(70)
- ✓ rotate-left(50)
- ✓ rotate-left(30)
- ✓ rotate-right(50)



6

Red-black tree operations

Insertion (overview)

• Steps

- ✓ insert the new node following standard BST rules
- ✓ color the new node **red** (to avoid violating the *root-to-null* rule)
- ✓ if parent is **black**, terminate (forms 3-node or 4-node)
- ✓ if parent is **red**, resolve the *red-red* violation using case-based rebalancing
- ✓ finally, ensure the root is **black**

• Violation resolution

- ✓ apply recoloring (preserves structure) and/or rotations (preserves order)
 - rotations are used when recoloring alone cannot restore properties

8

Insertion (detailed steps)

Initial insertion

- ✓ insert the new node following standard BST rules
- ✓ color the new node **red**
- ✓ apply the appropriate case resolution based on parent and uncle colors


Case 1: parent is **black**

- ✓ no *red-red* violation occurs
- ✓ all properties satisfied, terminate

Case 2: parent and uncle are both **red**

- ✓ recolor parent and uncle to black
- ✓ recolor grandparent to red
- ✓ recursively apply case analysis to grandparent

Example: case 2



consider 3 other analogous (sub)cases

<https://www.happycoders.eu/algorithms/red-black-tree-java/>

10

Insertion (detailed steps)

Case 3: triangle formation — parent is **red**, uncle is **black** (or null)

- ✓ rotate at parent to transform triangle into line formation
 - left rotation at parent if new node is right child of left parent
 - right rotation at parent if new node is left child of right parent
- ✓ proceed to case 4


Case 4: line formation — parent is **red**, uncle is **black** (or null)

- ✓ rotate at grandparent opposite to the line direction
 - right rotation at grandparent if line extends left
 - left rotation at grandparent if line extends right
- ✓ swap colors of parent and grandparent
- ✓ this case resolves the violation locally, no further propagation

Final step

- ✓ after all case resolutions, ensure the root node is colored **black**

Example: case 3 (triangle)





consider 1 other analogous (sub)case

<https://www.happycoders.eu/algorithms/red-black-tree-java/>

12

11

Example: case 4 (line)



consider 1 other analogous (sub)case

<https://www.happycoders.eu/algorithms/red-black-tree-java/>

13

Practice

- Insert the following keys into a RB-tree

✓ 10, 20, 30, 40, 50, 15, 25, 35, 45

Final remarks

Theoretical Equivalence

- due to the correspondence between red-black trees and 2-3-4 trees, the maximum height of a red-black tree with n nodes is $O(\log n)$, specifically at most $2 \log_2(n+1)$

Other operations

- search**: identical to standard BST search (colors are ignored, $O(\log n)$ time)
- delete**: follows BST deletion, then applies case-based rebalancing (not covered in lecture; more complex than insertion)

C++ Implementation

- STL containers `std::set` and `std::map` are typically implemented using red-black trees (though the standard doesn't mandate the specific data structure)

15

Analysis

Data Structure	Worst-case			Average-case			Ordered?
	insert at	delete	search	insert at	delete	search	
sequential (unordered)	$O(n)$	$O(n)$	$O(n)$	$O(n)$	$O(n)$	$O(n)$	No
sequential (ordered) binary search	$O(n)$	$O(n)$	$O(\log n)$	$O(n)$	$O(n)$	$O(\log n)$	Yes
BST	$O(n)$	$O(n)$	$O(n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	Yes
2-3-4	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	Yes
Red-Black	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	Yes

16

STL (ordered) containers

Ordered associative containers (STL)

Ordered associative containers implement sorted data structures that can be quickly searched – $O(\log n)$ complexity




Image credit: <https://www.scaler.com/topics/cpp/containers-in-cpp/>

18

What is the output? — set

```
#include <iostream>
#include <set>

int main() {
    std::set<int> numbers;

    numbers.insert(5);
    numbers.insert(2);
    numbers.insert(8);
    numbers.insert(2); // duplicate
    numbers.insert(5); // duplicate
    numbers.insert(2); // duplicate

    std::cout << "set elements: ";
    for (int num : numbers) {
        std::cout << num << " ";
    }
    std::cout << "\n";

    std::cout << "Count of 2 in set: " << numbers.count(2) << "\n";
}

return 0;
```

What is the output? — multiset

```
#include <iostream>
#include <set>

int main() {
    std::multiset<int> numbers;

    numbers.insert(5);
    numbers.insert(2);
    numbers.insert(8);
    numbers.insert(2); // duplicate
    numbers.insert(5); // duplicate
    numbers.insert(2); // duplicate

    std::cout << "multiset elements: ";
    for (int num : numbers) {
        std::cout << num << " ";
    }
    std::cout << "\n";

    std::cout << "Count of 2 in multiset: " << numbers.count(2) << "\n";
}

return 0;
```

19

20

What is the output? — map

```
#include <map>
#include <string>
#include <iostream>

int main() {
    std::map<std::string, int> freq_table;

    std::string words[] = {"apple", "banana", "apple", "cherry", "banana", "apple"};
    for (const auto& word : words) {
        freq_table[word]++;
        // creates entry with value 0 if key doesn't exist
    }

    if (freq_table.find("date") == freq_table.end()) {
        std::cout << "'date' not found in map\n";
    }

    std::cout << "Frequency of 'apple': " << freq_table["apple"] << "\n";
    std::cout << "Word frequencies (sorted by key):\n";
    for (const auto& pair : freq_table) {
        std::cout << pair.first << ": " << pair.second << "  ";
    }
    std::cout << "\n";
}

return 0;
}
```

21

What is the output? — multimap

```
#include <map>
#include <string>
#include <iostream>

int main() {
    std::multimap<std::string, std::string> phone_book;

    // Insert multiple numbers for same person
    phone_book.insert({"Alice", "555-1234"});
    phone_book.insert({"Alice", "555-5678"});
    phone_book.insert({"Bob", "555-9999"});
    phone_book.insert({"Alice", "555-0000"});

    std::cout << "Alice has " << phone_book.count("Alice") << " numbers\n";
    std::cout << "Phone Directory:\n";
    for (const auto& entry : phone_book) {
        std::cout << " " << entry.first << ": " << entry.second << "\n";
    }

    std::cout << "Alice's numbers:\n";
    auto range = phone_book.equal_range("Alice");
    for (auto it = range.first ; it != range.second ; ++it) {
        std::cout << " " << it->second << "\n";
    }

    return 0;
}
```

22